


Non sottovalutare la potenza dello Storage Account!!



Massimo Bonanni

Technical Trainer @ Microsoft

Sponsor & Org



DATA SKILLS
UNDERSTANDING THE WORLD




Lucient¹
ITALIA



A day like many others, in a company like many others...

Carlo
Project
Manager



Maria
Solution
Architect



The requirements

- ✓ We produce an average of 100 Mb of new contents every hour in two different locations
- ✓ Contents are frequently accessed in the first 10 days, very rarely after a month, and can be deleted after 1 year
- ✓ The contents cannot be deleted for 10 days after publication
- ✓ The maximum downtime in a single month must be under 5 minutes
- ✓ The costs must be kept to a minimum
- ✓ We need daily statistics on the number of contents available, and on the space occupied



**What if we used the
Storage Account?**

?????

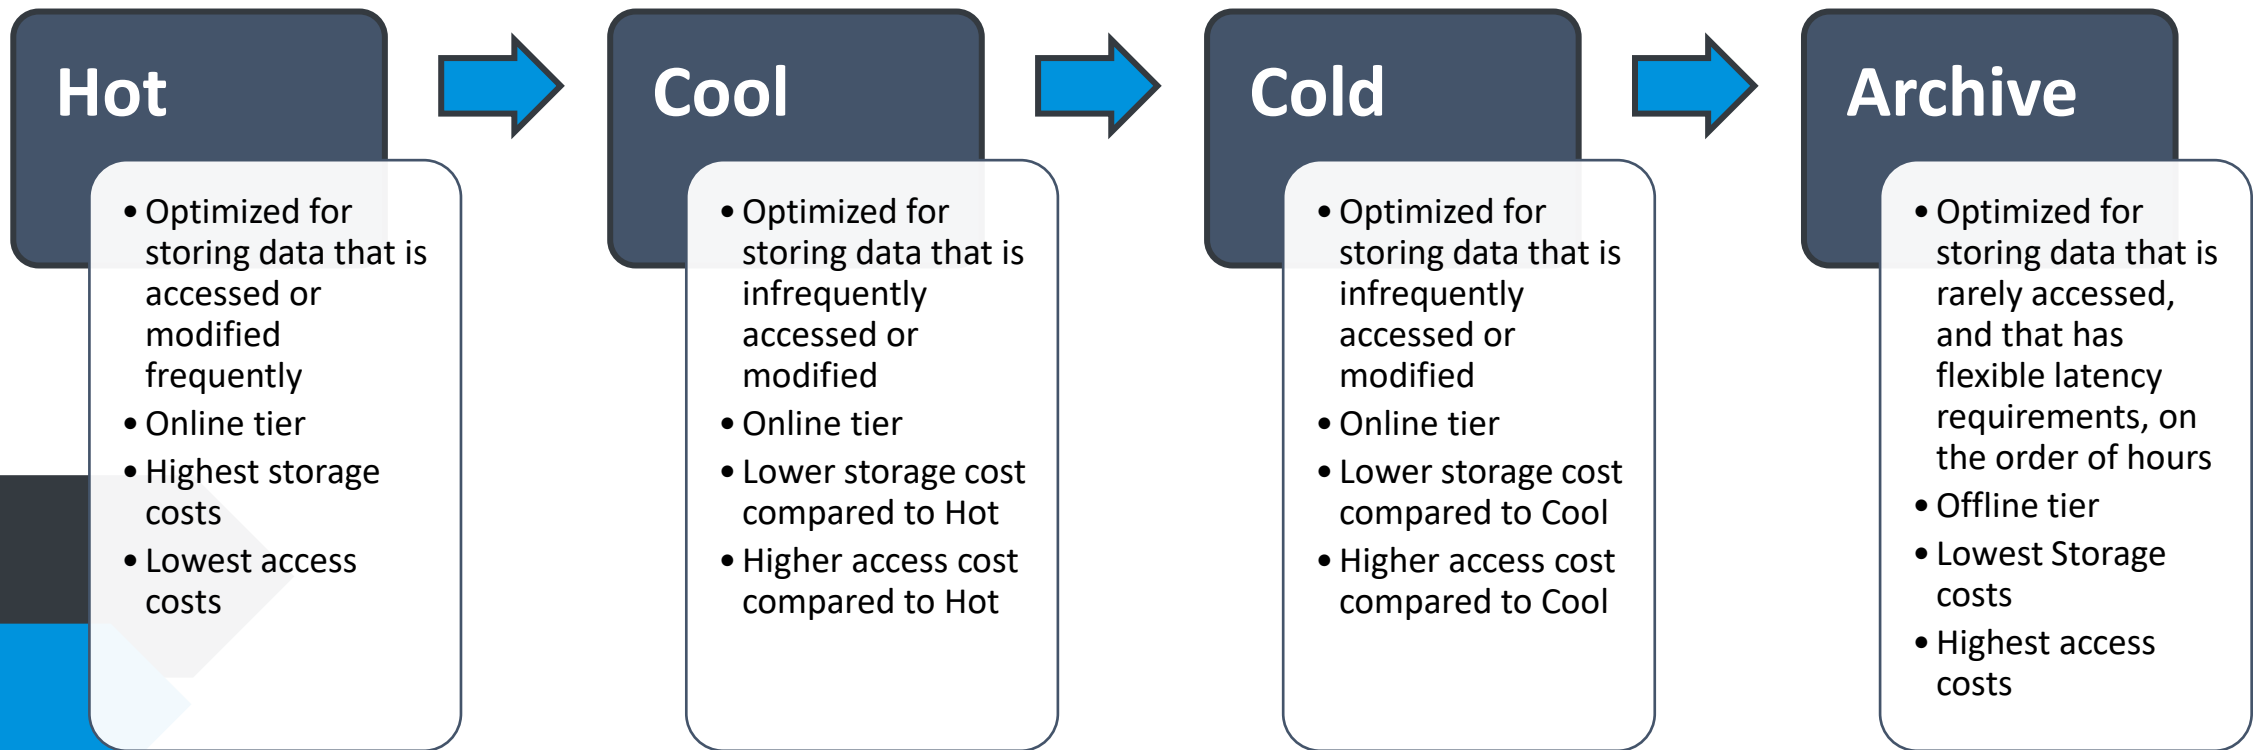


How do we minimize costs?



Access Tiers

Azure storage offers different access tiers so that you can store your blob data in the most cost-effective manner based on how it's being used.



Access Tiers

Contents will be stored in the **Hot** tier for the first **10 days**, then will be "moved" to the **Cool** tier for the next **30 days**, and then will be "moved" to the **Cold** tier. We can finally move them after 1 year to the **Archive** tier (or delete).

- Highest storage costs
- Lowest access costs

- Lower storage cost compared to Hot
- Higher access cost compared to Hot


- Lower storage cost compared to Cool
- Higher access cost compared to Cool

Archive

- Optimized for storing rarely accessed data

flexible retention requirements on the order of years

- Lowest storage costs
- Highest access costs



But... who "moves" the contents?



Blob Lifecycle Management




Azure Storage lifecycle management is a **rule-based policy** that allows you to “move” blob data to the appropriate access tiers or to expire data at the end of the data lifecycle.

A lifecycle management policy is comprised of one or more rules that define a set of actions to take based on a condition being met.

Blob Lifecycle Management

Let's create a rule that moves contents to the **Cool** tier **10 days** after their creation and moves them to the **Cold** tier after **30 days**...and deletes them after **1 year**.

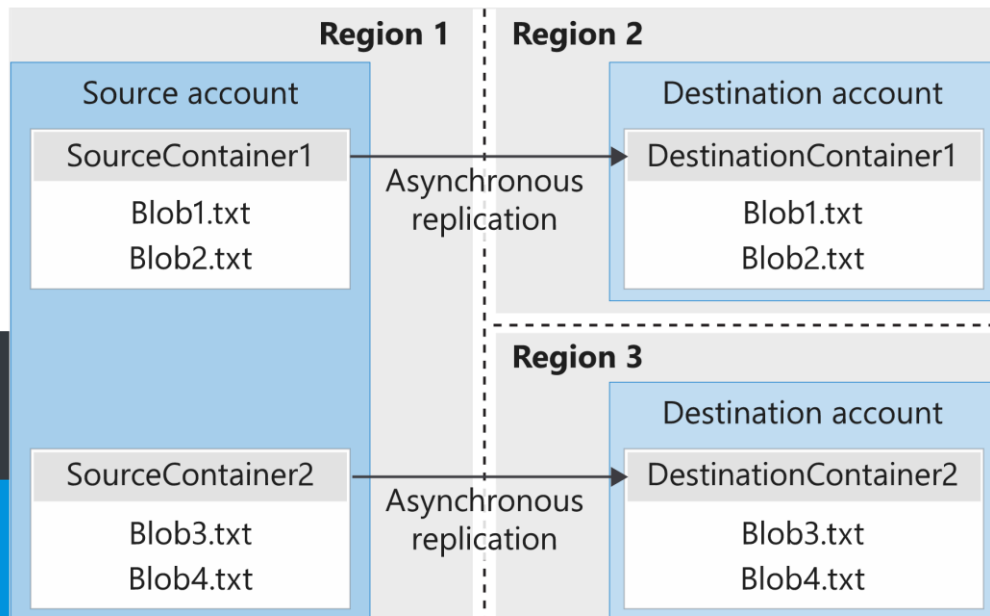


But...we have **multiple storages in multiple regions** and all our users must see the same contents in all the regions.



Object Replication

Object replication asynchronously copies block blobs between a source storage account and a destination account.



Minimizing latency



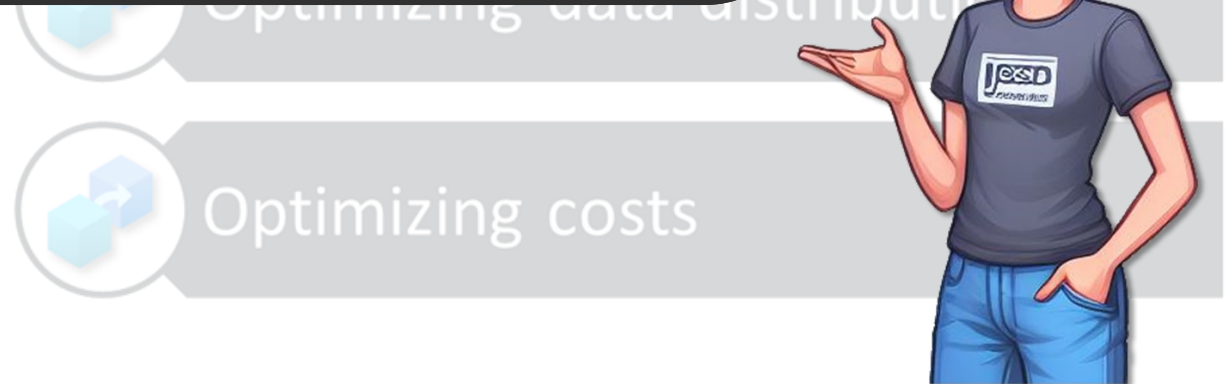
Optimizing data distribution



Optimizing costs

Object Replication

Contents produced in one region will be **replicated to the other regions**, so the web apps in each region will read all contents **locally**.



Cool...but...we need also to create a software that analyzes, **every day**, the contents in each storage to save **statistics**.



Storage Inventory

Blob inventory feature provides an overview of your containers, blobs, snapshots, and blob versions within a storage account managed by Azure.



Inventory reports for blobs and containers



Custom Schema



CSV and Apache Parquet output format



Manifest file and Azure Event Grid event per inventory report

Storage Inventory

Blob inventory feature provides an overview of your containers, blobs, snapshots, and blob versions within a storage account managed by Azure.

Using inventory, **every day**, we will have a **CSV file** with all the info we need to create our statistics.

We will use **Event Grid and Azure Functions** to extract data from Inventory file when it will be generated by Azure.



Manifest file and Azure Event Grid event per inventory refresh



...and what about **avoid** contents
deletion in the first **10 days** after
their creation?



Immutable Storage

Immutable storage for Azure Blob Storage enables users to store business-critical data in a **WORM (Write Once, Read Many)** state.

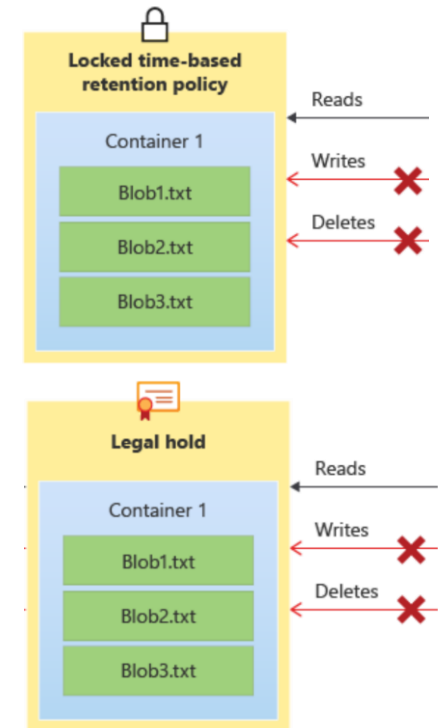
While in a WORM state, data cannot be modified or deleted for a user-specified interval.



Time-based retention policies



Legal hold policies



Immutable Storage

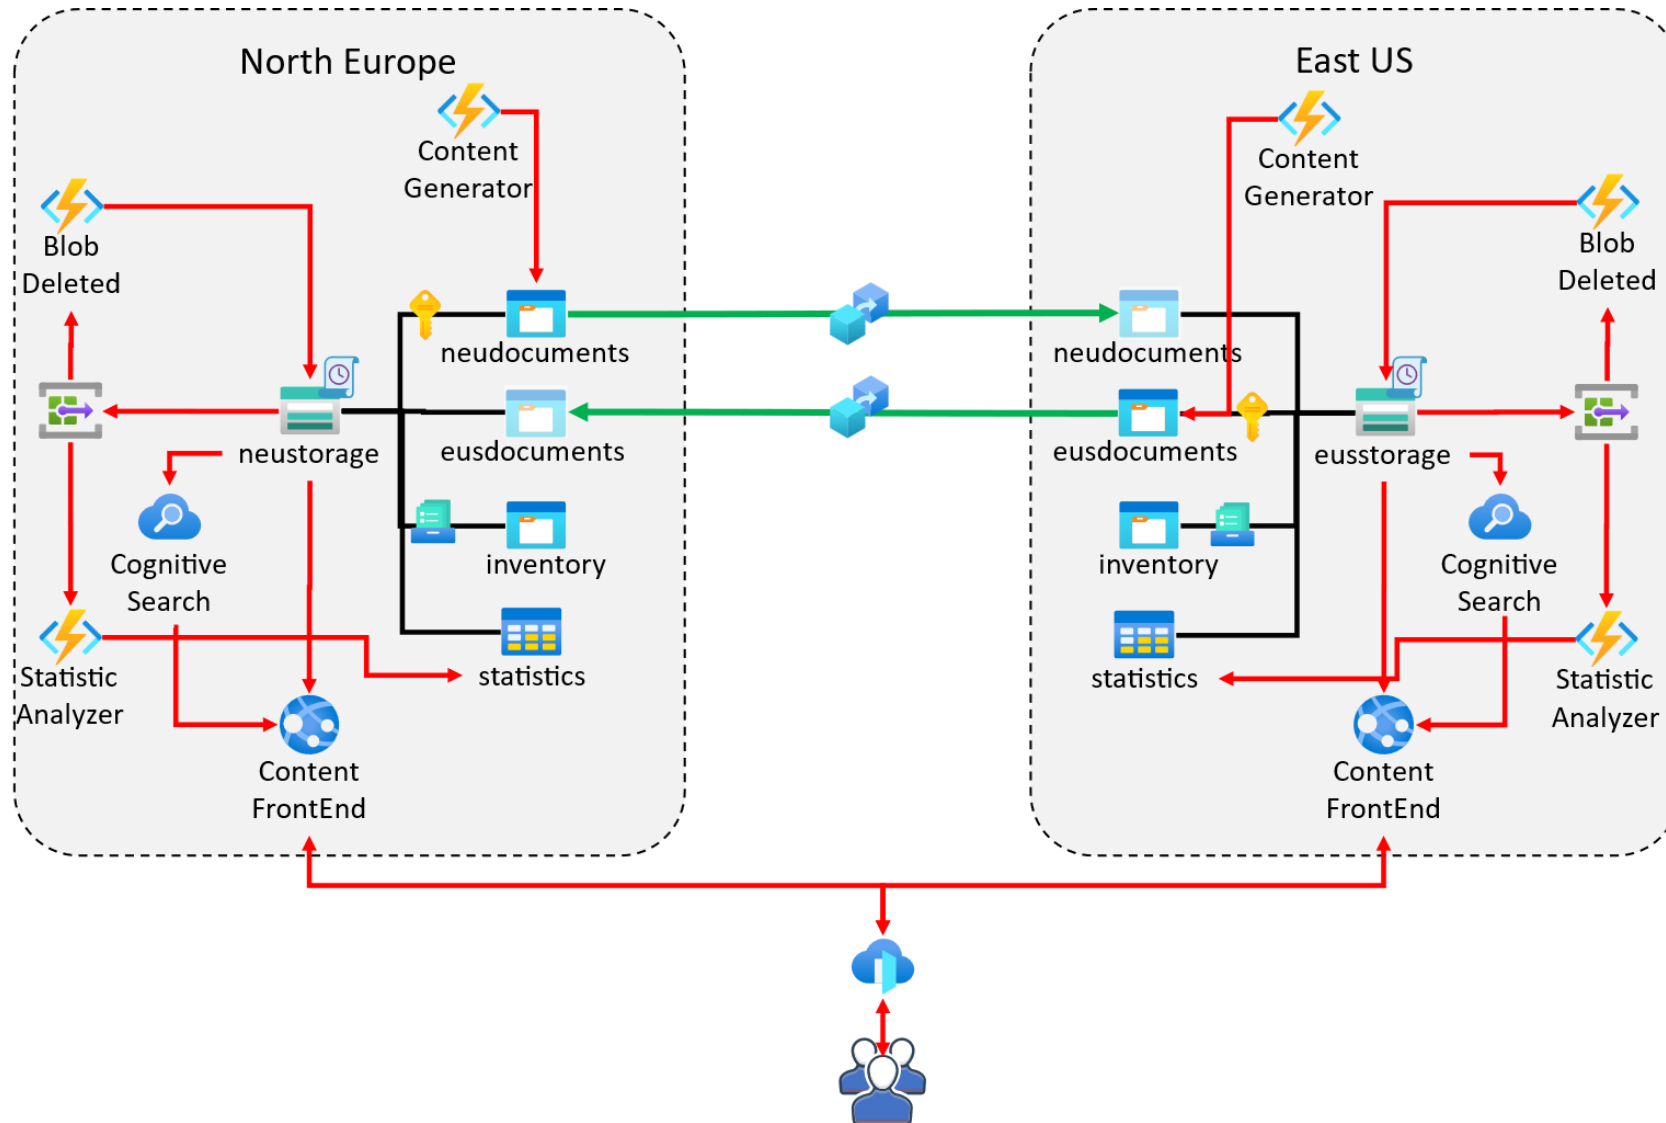
Immutable storage for Azure Blob Storage enables users to store business-critical data.

We will set a **time-based policy** to avoid deletion in the first **10 days** after content creation.

Legal hold policies



The solution



DEMO



Storage Content Platform



[massimobonanni/StorageContentPlatform](https://github.com/massimobonanni/StorageContentPlatform)

...are you sure we can provide at most 5 minutes of downtime every month?



The SLA of the solution

5 minutes of downtime every month

➔ **SLA 99.9885 %**

SLA 1 region

99.95%

AppService

99.99%

Storage
(RA-GRS)

$$SLA_{sr} = 99.95\% * 99.99\% \\ = 99.94\%$$

SLA 2 regions

Region1

99.94%

Region2

99.94%

$$SLA_{mr} = 100\% - (0.06\% * 0.06\%) \\ \cong 99.9999\%$$

SLA full solution

99.99%

FD

99.9999%



Multi
Region

$$SLA_{tot} = 99.99\% * SLA_{mr} \\ \cong \mathbf{99.9899\%}$$







...and how will cost it?



Costs




Feature-to-Requirement mapping

	Contents in different regions	Object Replication
	Contents are accessed differently depending on the age	Access Tiers
	The contents cannot be deleted for 10 days after publication	Time-based retention policy
	The maximum downtime must be under 5 minutes	Multi-region
	The costs must be kept to a minimum	Lifecycle rules (access tiers) + Serverless
	Daily statistics	Inventory


**Why don't we use
Storage Account?**





20% discount code

AZURECOOKBOOK24



Massimo Bonanni

Technical Trainer @ Microsoft

massimo.bonanni@microsoft.com

@massimobonanni





#64 PARMA 2024

Thanks!!!

